

# Numerical Optimization

Lecture Notes #12

Conjugate Gradient Methods — Nonlinear CG

Fall 2019

## Outline

### 1 Linear Conjugate Gradient Methods

- Recap
- Conjugate Gradient Algorithms
- The Effect of Preconditioning — CG vs. PCG(M)

### 2 Nonlinear Conjugate Gradient Methods

- New Ideas... Fletcher-Reeves, etc...
- Practical Considerations
- Convergence

### 3 Projects

- Separate Handouts, etc...

## Quick Recap: Linear Conjugate Gradient Methods

We have introduced the Conjugate Gradient (CG) and Preconditioned Conjugate Gradient (PCG) methods for solution of the linear system  $A\bar{x} = \bar{b}$ , where  $A$  is symmetric positive definite.

**Linear CG** is guaranteed to converge in  $n$  iterations, but as we have seen, in many cases — eigenvalue clustering and/or  $r < n$  distinct eigenvalues, convergence is much faster.

We briefly discussed preconditioning, where we use a simplified version  $M \approx A$ , and hope that  $M^{-1}A \approx I$  has a favorable eigenvalue spectrum. We must be able to solve  $M\bar{y} = \bar{r}$  fast.

### Today:

- (i) An example of CG vs. PCG performance.
- (ii) Non-linear CG.
- (iii) Projects!

# The CG Algorithm (version 1.0, "Standard")

## Algorithm: Conjugate Gradient

Given  $A$ ,  $\bar{\mathbf{b}}$  and  $\bar{\mathbf{x}}_0$ :

$$\bar{\mathbf{r}}_0 = A\bar{\mathbf{x}}_0 - \bar{\mathbf{b}}, \quad \bar{\mathbf{p}}_0 = -\bar{\mathbf{r}}_0, \quad k = 0$$

while (  $\|\mathbf{r}_k\| > 0$ , or other stopping condition )

$$\alpha_k = \frac{\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k}, \quad \begin{array}{l} \text{Store the vector } A\bar{\mathbf{p}}_k \\ \text{and the scalar } \bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k \end{array}$$

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{r}}_{k+1} = \bar{\mathbf{r}}_k + \alpha_k A \bar{\mathbf{p}}_k$$

$$\beta_{k+1} = \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{r}}_{k+1}}{\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k}, \quad \text{Keep numerator for next step!}$$

$$\bar{\mathbf{p}}_{k+1} = -\bar{\mathbf{r}}_{k+1} + \beta_{k+1} \bar{\mathbf{p}}_k$$

$$k = k + 1$$

end-while

## Preconditioned CG Algorithm (a.k.a. "PCG(M)")

### Algorithm: PCG

Given  $A$ ,  $\mathbf{M} = \mathbf{C}^T \mathbf{C}$ ,  $\bar{\mathbf{b}}$  and  $\bar{\mathbf{x}}_0$ : compute  $\bar{\mathbf{r}}_0 = A\bar{\mathbf{x}}_0 - \bar{\mathbf{b}}$ ,  
 $\bar{\mathbf{y}}_0 = \mathbf{M}^{-1}\bar{\mathbf{r}}_0$ ,  $\bar{\mathbf{p}}_0 = -\bar{\mathbf{y}}_0$ ,  $k = 0$

while (  $\|r_k\| > 0$ , or other stopping condition )

$$\alpha_k = \frac{\bar{\mathbf{r}}_k^T \bar{\mathbf{y}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k}, \quad \begin{array}{l} \text{Store the vector } A\bar{\mathbf{p}}_k \\ \text{and the scalar } \bar{\mathbf{r}}_k^T \bar{\mathbf{y}}_k \end{array}$$

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{r}}_{k+1} = \bar{\mathbf{r}}_k + \alpha_k A \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{y}}_{k+1} = \mathbf{M}^{-1} \bar{\mathbf{r}}_{k+1}$$

$$\beta_{k+1} = \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{y}}_{k+1}}{\bar{\mathbf{r}}_k^T \bar{\mathbf{y}}_k}, \quad \begin{array}{l} \text{Save the numerator for next step!} \end{array}$$

$$\bar{\mathbf{p}}_{k+1} = -\bar{\mathbf{y}}_{k+1} + \beta_{k+1} \bar{\mathbf{p}}_k$$

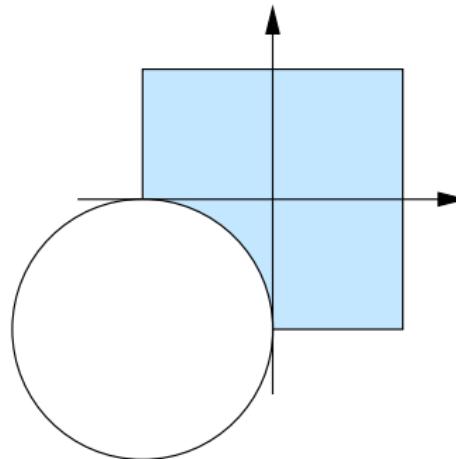
$$k = k + 1$$

end-while

## Example: CG vs. PCG(M) Performance

1 of 7

**Problem:** Solve  $\nabla^2 u(x, y) = f(x, y)$  in the domain

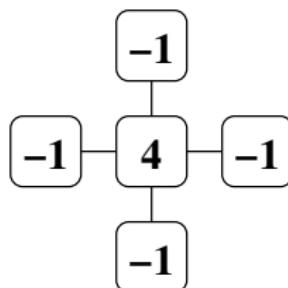


$$D = \{(x, y) : -1 \leq x, y \leq 1\} - \{(x, y) : (x + 1)^2 + (y + 1)^2 < 1\}$$

Set  $u(x, y) = 0$  on  $\Gamma = \partial D$  (Dirichlet Boundary Conditions).

## Example: CG vs. PCG(M) Performance

2 of 7



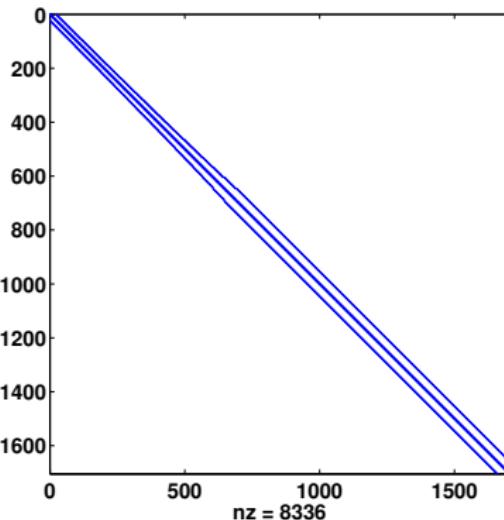
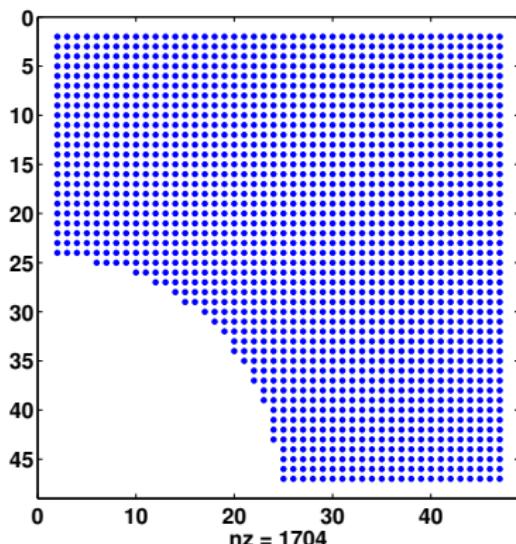
**Figure:** We discretize  $\nabla^2 u(x, y)$  by the standard 5-point finite difference approximation of the Laplacian.

We study the numerical solution of the resulting linear system  $A\bar{\mathbf{u}} = \bar{\mathbf{f}}$  for varying discretizations of the square (from  $2 \times 2$  to  $64 \times 64$  grids.)

We look at CG, PCG(M) with  $M$  being the tri-diagonal preconditioner, and PCG(M) with  $M = \tilde{L}\tilde{L}^T$ , where  $\tilde{L}$  is given by the incomplete (zero fill-in) Cholesky factorization.

## Example: CG vs. PCG(M) Performance

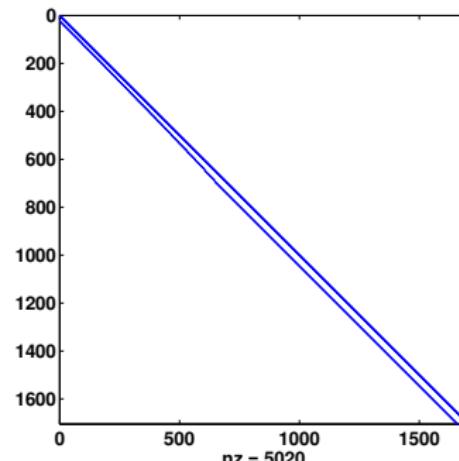
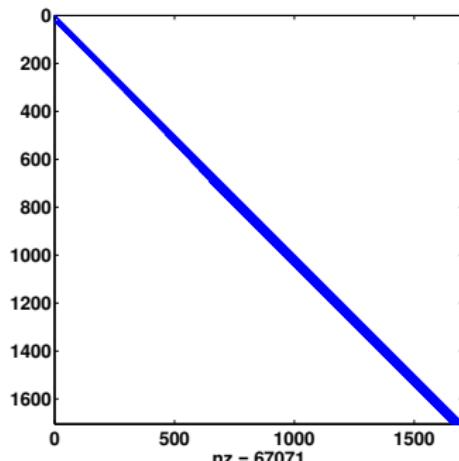
3 of 7



**Figure:** We discretize  $\nabla^2 u(x, y)$  by the standard 5-point finite difference approximation of the Laplacian on the numerical domain (illustrated on the left with a  $48 \times 48$  grid), the corresponding matrix  $A$  is illustrated to the right; it has a tri-diagonal component, and two additional elements on every row — the bandwidth is not constant due to the cut-out in the domain.

## Example: CG vs. PCG(M) Performance

4 of 7



**Figure:** To the left we see the  $L$  given by complete Cholesky factorization — we notice how the entire band fills in, we get a total of 67,071 non-zero entries. To the right we see the  $\tilde{L}$  given by **incomplete Cholesky factorization** — here we only get 5,020 non-zero entries. ( $A$  had 8,336 non-zero entries)

We will use the preconditioners  $M = \tilde{L}\tilde{L}^T$ , and  $M = \text{tridiag}(A)$ .

## Example: CG vs. PCG(M) Performance

5 of 7

$n_{\text{GRID}}$	$48^2$	$64^2$
$A$	$1704 \times 1704$	$3094 \times 3094$

 $L = \text{ichol}(A), LL^T \approx A$ 

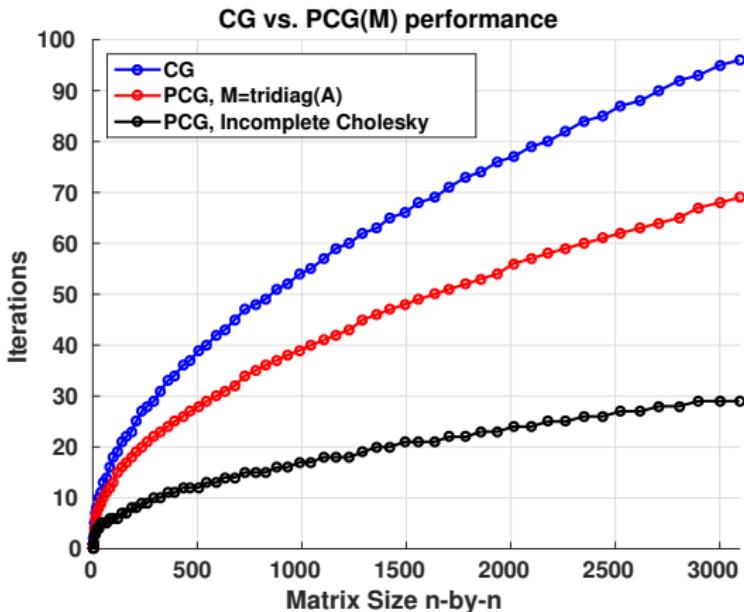
$\ A - LL^T\ _F / \ A\ _F$	0.0900	0.0907
$\text{cond}(A)$	924.5193	1656.936
$\text{cond}(L^{-1}AL^{-T})$	133.4733	238.4772

 $L = \text{chol}(\text{tridiag}(A)), LL^T \approx A$ 

$\ A - LL^T\ _F / \ A\ _F$	0.3131	0.3139
$\text{cond}(A)$	924.5193	1656.936
$\text{cond}(L^{-1}AL^{-T})$	463.1769	829.6574

## Example: CG vs. PCG(M) Performance

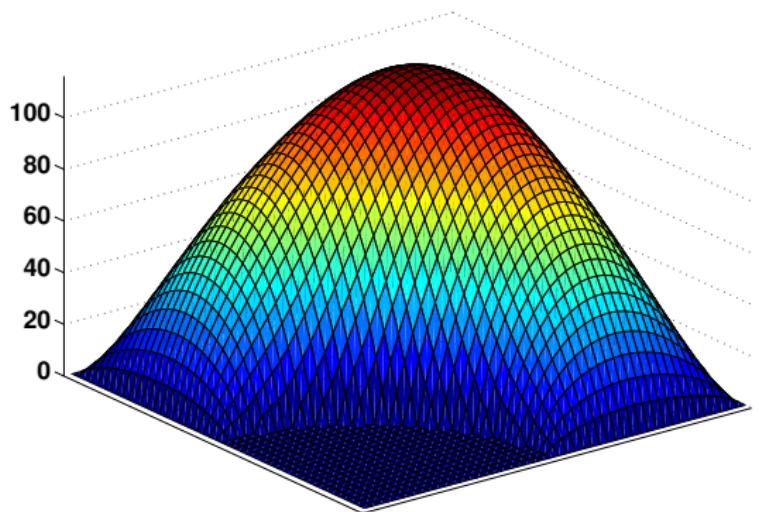
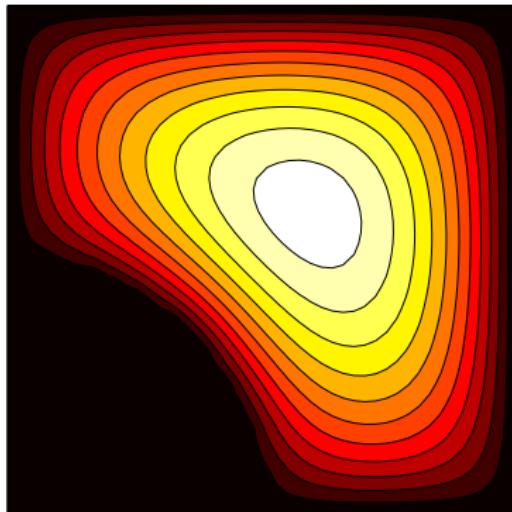
6 of 7



**Figure:** The performance of **CG** and **PCG(M)** on our test problem. The discretization of the square  $[-1, 1]^2$  ranges from  $2 \times 2$  to  $64 \times 64$ , which gives us a matrix  $A$  of dimensions ranging from  $3 \times 3$  to  $3094 \times 3094$ . The stopping criteria was a relative reduction of  $\|\bar{r}\|$  by  $10^{-6}$ .

## Example: CG vs. PCG(M) Performance

7 of 7



**Figure:** The solution to our test problem on the  $48 \times 48$  grid, with the right-hand-side  $f(x, y) = 1$ .

## Nonlinear Conjugate Gradient Methods

We now turn our attention to making the CG methods useful for optimization problems (the non-linear situation).

The **Fletcher-Reeves** (CG-FR, published in 1964) extension requires two modifications to the CG algorithm:

- 1: The computation of the step length  $\alpha_k$  is replaced by a line-search which minimizes the non-linear objective  $f(\cdot)$  along the search direction  $\bar{p}_k$ .
- 2: The instances of the residual  $\bar{r}$  (which are just  $\nabla\Phi(\cdot)$  for the quadratic objective in standard CG) are replaced by the gradient of the non-linear objective  $\nabla f(\cdot)$ .

---

Fletcher, R., and Reeves, C. M. "Function minimization by conjugate gradients." *The computer journal*, 7, no. 2 (1964), 149-154.

## The Fletcher-Reeves FR-CG Algorithm

### Algorithm: Fletcher-Reeves

Given  $\bar{\mathbf{x}}_0$ :

Evaluate  $f_0 = f(\bar{\mathbf{x}}_0)$ ,  $\nabla f_0 = \nabla f(\bar{\mathbf{x}}_0)$ .

Set  $\bar{\mathbf{p}}_0 = -\nabla f_0$ ,  $k = 0$

while (  $\|\nabla f_k\| > 0$ , ... )

$\alpha_k = \text{linesearch}(\dots)$

$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k$

$\nabla f_{k+1} = \text{Evaluate } \nabla f(\bar{\mathbf{x}}_{k+1})$

$\beta_{k+1}^{\text{FR}} = \frac{\nabla f_{k+1}^T \nabla f_{k+1}}{\nabla f_k^T \nabla f_k}$ , Save numerator!

$\bar{\mathbf{p}}_{k+1} = -\nabla f_{k+1} + \beta_{k+1}^{\text{FR}} \bar{\mathbf{p}}_k$

$k = k + 1$

end-while

## Comments: The Fletcher-Reeves FR-CG Algorithm

**Sanity check:** If  $f(\bar{\mathbf{x}})$  is a strongly convex quadratic, and  $\alpha_k$  the exact minimizer, then FR-CG reduces to linear CG.

Each iteration requires evaluation of the objective function (for the line-search), and the gradient of the objective. — No Hessian evaluation, nor matrix operations are required. **Good** for large non-linear optimization problems.

If we require that  $\alpha_k$  satisfies the **strong Wolfe conditions**

$$\begin{aligned} f(\bar{\mathbf{x}}_k + \alpha \bar{\mathbf{p}}_k) &\leq f(\bar{\mathbf{x}}_k) + c_1 \alpha \bar{\mathbf{p}}_k^T \nabla f_k \\ |\bar{\mathbf{p}}_k^T \nabla f(\bar{\mathbf{x}}_k + \alpha \bar{\mathbf{p}}_k)| &\leq c_2 |\bar{\mathbf{p}}_k^T \nabla f_k| \end{aligned}$$

where  $0 < c_1 < c_2 < \frac{1}{2}$ , then FR-CG converges globally.

## Variants: The Polak-Ribi  re (PR-CG) Method

1 of 2

The following modification to FR-CG was suggested by Polak-Ribi  re

$$\beta_{k+1}^{\text{FR}} = \frac{\nabla f_{k+1}^T \nabla f_{k+1}}{\nabla f_k^T \nabla f_k} \quad \rightarrow \quad \beta_{k+1}^{\text{PR}} = \frac{\nabla f_{k+1}^T (\nabla f_{k+1} - \nabla f_k)}{\nabla f_k^T \nabla f_k}$$

when  $f$  is a strongly convex quadratic, and the line search is exact, the gradients are orthogonal and  $\beta_{k+1}^{\text{FR}} = \beta_{k+1}^{\text{PR}}$ .

On general non-linear objectives, an inexact line-searches PR-CG tends to be **more robust** and **more efficient** than FR-CG.

---

Polak, Elijah, and Gerard Ribiere. "Note sur la convergence de m  thodes de directions conjugu  s." *Revue fran  aise d'informatique et de recherche op  rationnelle*. S  rie rouge 3, no. 16 (1969): 35-43.

## Variants: The Polak-Ribière (PR-CG) Method

2 of 2

**One problem:** The strong Wolfe conditions **do not** guarantee that  $\bar{p}_k$  is always descent direction for PR-CG. In order to fix this,  $\beta$  is defined to be

$$\beta_{k+1}^+ = \max(\beta_{k+1}^{\text{PR}}, 0)$$

the resulting algorithm is known as PR+.

There are a number of other choices for  $\beta$  in the literature, but they are not (in general) more efficient than Polak-Ribière PR-CG/PR+.

## Practical Considerations

1 of 2

If the line-search uses quadratic (or cubic) interpolation along the search direction  $\bar{p}_k$ , then if/when  $f(\cdot)$  is a strictly convex quadratic, the step lengths  $\alpha_k$  will be the exact 1D-minimizers  $\Rightarrow$  the non-linear algorithm reduces to linear CG. [THIS IS HIGHLY DESIRABLE!]

**Restarting:** CG gets its favorable convergence properties from the conjugacy of the search directions **near** the optimum. If we start “far” from the optimum, the algorithm does not necessarily gain anything from maintaining this conjugacy.

Therefore, we should periodically restart the algorithm, by setting  $\beta = 0$  (i.e. taking a steepest-descent step).

The  $n$ -step convergence is only guaranteed when we start with a steepest-descent step, and the model is quadratic. Hence a restart close to  $\bar{x}^*$  will (approximately) guarantee  $n$ -step convergence.

## Practical Considerations: Restarting Conditions

2 of 2

**Restarting conditions:** The most common condition is based on the fact that for the strictly quadratic objective, the residuals are orthogonal. Hence, when two consecutive residuals are “far” from orthogonal

$$\frac{\nabla f_k^T \nabla f_{k-1}}{\nabla f_k^T \nabla f_k} \geq \nu \sim 0.1$$

a restart is triggered.

The formula

$$\beta_{k+1}^+ = \max(\beta_{k+1}^{\text{PR}}, 0)$$

in PR+ can be viewed as a restart-condition. This is not practical since these “restarts” are very infrequent — in practice  $\beta_{k+1}^{\text{PR}}$  is positive most of the time.

## Nonlinear CG: Global Convergence

**Linear CG:** Global convergence properties well understood, and optimal.

**Nonlinear CG:** Convergence properties not so well understood, except in special cases. The behavior is sometimes surprising and bizarre!

We look at some results, under the following non-restrictive assumptions

### Assumptions:

- (i) The level set  $\mathcal{L} = \{\bar{\mathbf{x}} \in \mathbb{R}^n : f(\bar{\mathbf{x}}) \leq f(\bar{\mathbf{x}}_0)\}$  is bounded.
- (ii) In some neighborhood  $\mathcal{N}$  of  $\mathcal{L}$ , the objective function  $f$  is Lipschitz continuously differentiable, i.e. there exists a constant  $L > 0$  such that

$$\|\nabla f(\bar{\mathbf{x}}) - \nabla f(\bar{\mathbf{y}})\| \leq L\|\bar{\mathbf{x}} - \bar{\mathbf{y}}\|, \quad \forall \bar{\mathbf{x}}, \bar{\mathbf{y}} \in \mathcal{N}$$

## Global Convergence: FR-CG

## Theorem

Suppose that the assumptions hold, and that FR-CG is implemented with a line search which satisfies the strong Wolfe conditions, with  $0 < c_1 < c_2 < \frac{1}{2}$ . Then

$$\liminf_{k \rightarrow \infty} \|\nabla f_k\| = 0.$$

This does not say that the limit of the sequence of gradients  $\{\nabla f_k\}$  is zero; but it does tell us that at least the sequence is not bounded away from zero.

If, however, we restart the algorithm every  $n$  steps, we get  $n$ -step quadratic convergence:

$$\|\bar{x}_{k+n} - \bar{x}^*\| = \mathcal{O}(\|\bar{x}_k - \bar{x}^*\|^2).$$

## Global Convergence: PR-CG

In practice PR-CG performs better than FR-CG, but we cannot prove a theorem like the one for FR-CG on the previous slide.

The following surprising result **can** be shown:

### Theorem

*Consider the Polak-Ribiere PR-CG method with an ideal line search. There exists a twice continuously differentiable objective function  $f : \mathbb{R}^3 \rightarrow \mathbb{R}$  and a starting point  $\bar{x}_0 \in \mathbb{R}^3$  such that the sequence of gradients  $\{\|\nabla f_k\|\}$  is bounded away from zero.*

### The modification (PR+)

$$\beta_{k+1}^+ = \max(\beta_{k+1}^{\text{PR}}, 0)$$

fixes this strange behavior, and it is possible to show global convergence for PR+.



# T.P.S REPORT

## CO V E R S H E E T

Prepared By: \_\_\_\_\_ Date: \_\_\_\_\_

System: \_\_\_\_\_ Program Language: \_\_\_\_\_ Platform: \_\_\_\_\_ OS: \_\_\_\_\_

Unit Code: \_\_\_\_\_ Customer: \_\_\_\_\_

Unit Code Tested: \_\_\_\_\_

Due Date: \_\_\_\_\_ Approved By: \_\_\_\_\_

Test Date: \_\_\_\_\_ Tested By: \_\_\_\_\_

Total Run Time: \_\_\_\_\_ Total Error Count: \_\_\_\_\_

Error Reference: \_\_\_\_\_

Errors Logged: \_\_\_\_\_ Log Location: \_\_\_\_\_

Passed: \_\_\_\_\_ Moved to Production: \_\_\_\_\_

Comments: \_\_\_\_\_  
\_\_\_\_\_  
\_\_\_\_\_

C O N F I D E N T I A L

# Index